

# AI Applications Lecture 18

Parameter Efficient Fine Tuning and LoRA

---

SUZUKI, Atsushi  
Jing WANG

# Outline

Introduction

Preparation: Mathematical Notations Revisited

The Idea of PEFT: Representing Model Differences with Low-Dimensional Parameters

The Idea of Low Rank Adaptation (LoRA)

Specific Application of LoRA: Example with Stable Diffusion 1.5

Summary and Next Lecture Preview

# Introduction

---

## Introduction (1)

The neural networks currently used in practice in the field of generative AI have an enormous number of parameters. Since machine learning is a framework for determining the parameters of a parametric function from training data, an amount of training data corresponding to the number of parameters to be determined is required.

## Introduction (1)

The neural networks currently used in practice in the field of generative AI have an enormous number of parameters. Since machine learning is a framework for determining the parameters of a parametric function from training data, an amount of training data corresponding to the number of parameters to be determined is required.

For example, if one attempts to minimize an objective function using a gradient method, the objective function itself must contain enough information to sufficiently judge the quality of the parameters for the task one wants to solve. This requires an amount of training data corresponding to the number of parameters.

## Introduction (2)

Collecting a large amount of training data is accompanied by difficulties, and even if it is collected, learning on it requires a large amount of computation. Therefore, if possible (i.e., if there are no major performance problems), one would like to reduce the amount of training data used. For learning to succeed even with less training data, the number of parameters changed during learning needs to be small.

## Introduction (2)

Collecting a large amount of training data is accompanied by difficulties, and even if it is collected, learning on it requires a large amount of computation. Therefore, if possible (i.e., if there are no major performance problems), one would like to reduce the amount of training data used. For learning to succeed even with less training data, the number of parameters changed during learning needs to be small.

For the reasons mentioned above, **one would like to reduce the number of parameters changed during learning, if possible**. Of course, in a situation where one is building a model by training from scratch, if the task itself is complex and the desired input-output relationship is complicated, a high-dimensional checkpoint (a large number of parameters) and a large-scale architecture are necessary to represent that complex input-output relationship, and data for training it is also necessary. Thus, the policy of reducing the number of parameters to be learned is inherently unreasonable

## Introduction (3)

On the other hand, if one is in a situation where one **already has** a model  $f_\theta$  that gives a good input-output relationship, and it is expected that a better input-output relationship can be obtained by modifying it slightly (in some sense), then a scheme of **learning the minimum number of parameters just enough to represent the part corresponding to that slight modification** might be practically feasible. The idea of starting from an existing model and modifying it slightly to obtain a better model is called **fine-tuning**, and the idea of realizing this by learning a small number of parameters is called **PEFT (parameter-efficient fine-tuning)**. In this lecture, we will learn about **LoRA (low-rank adaptation)** as a typical example of PEFT [3, 2].

## Learning Outcomes

After completing this lecture, students should be able to do the following:

- Control the behavior of generative AI using LoRA.
- Explain the advantages of LoRA.
- Distinguish and explain what LoRA does mathematically and what it does in implementation.

# **Preparation: Mathematical Notations Revisited**

---

## Mathematical Notations (1)

- **Definition:**

- (LHS)  $\coloneqq$  (RHS): Indicates that the left-hand side is defined by the right-hand side. For example,  $a \coloneqq b$  indicates that  $a$  is defined by  $b$ .

## Mathematical Notations (2): Sets

- **Set (Set):**

- Sets are often denoted by uppercase calligraphic letters. Example:  $\mathcal{A}$ .
- $x \in \mathcal{A}$ : Indicates that the element  $x$  belongs to the set  $\mathcal{A}$ .
- $\{\}$ : Empty set.
- $\{a, b, c\}$ : The set consisting of elements  $a, b, c$  (extensional notation).
- $\{x \in \mathcal{A} \mid P(x)\}$ : The set of elements in set  $\mathcal{A}$  for which the proposition  $P(x)$  is true (intensional notation).
- $\mathbb{R}$ : The set of all real numbers.
- $\mathbb{R}_{>0}$ : The set of all positive real numbers.
- $\mathbb{R}_{\geq 0}$ : The set of all non-negative real numbers.
- $\mathbb{Z}$ : The set of all integers.
- $\mathbb{Z}_{>0}$ : The set of all positive integers.
- $\mathbb{Z}_{\geq 0}$ : The set of all non-negative integers.
- $[1, k]_{\mathbb{Z}} := \{1, 2, \dots, k\}$ : For a positive integer  $k$ , the set of integers from 1 to  $k$ .

## Mathematical Notations (3): Functions

- **Function (Function):**

- $f : \mathcal{X} \rightarrow \mathcal{Y}$ : Indicates that the function  $f$  is a map that takes elements from set  $\mathcal{X}$  as input and outputs elements from set  $\mathcal{Y}$ .
- $y = f(x)$ : Indicates that the output is  $y \in \mathcal{Y}$  when  $x \in \mathcal{X}$  is input to the function  $f$ .

## Mathematical Notations (4): Vectors

- **Vector (Vector):**

- In this course, a vector refers to a column of numbers.
- Vectors are denoted by bold italic lowercase letters. Example:  $v$ .
- $v \in \mathbb{R}^n$ : Indicates that the vector  $v$  is an  $n$ -dimensional real vector.
- The  $i$ -th element of vector  $v$  is denoted as  $v_i$ .

$$v = \begin{bmatrix} v_1 \\ v_2 \\ \vdots \\ v_n \end{bmatrix}. \quad (1)$$

## Mathematical Notations (5): Matrices

- **Matrix (Matrix):**

- Matrices are denoted by bold italic uppercase letters. Example:  $\mathbf{A}$ .
- $\mathbf{A} \in \mathbb{R}^{m,n}$ : Indicates that the matrix  $\mathbf{A}$  is an  $m \times n$  real matrix.
- The element in the  $i$ -th row and  $j$ -th column of matrix  $\mathbf{A}$  is denoted as  $a_{i,j}$ .

$$\mathbf{A} = \begin{bmatrix} a_{1,1} & a_{1,2} & \cdots & a_{1,n} \\ a_{2,1} & a_{2,2} & \cdots & a_{2,n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m,1} & a_{m,2} & \cdots & a_{m,n} \end{bmatrix}. \quad (2)$$

## Mathematical Notations (6): Transpose

- **Matrix (Matrix), continued:**

- The transpose of matrix  $A$  is denoted as  $A^\top$ . If  $A \in \mathbb{R}^{m,n}$ , then  $A^\top \in \mathbb{R}^{n,m}$ , and

$$A^\top = \begin{bmatrix} a_{1,1} & a_{2,1} & \cdots & a_{m,1} \\ a_{1,2} & a_{2,2} & \cdots & a_{m,2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1,n} & a_{2,n} & \cdots & a_{m,n} \end{bmatrix} \quad (3)$$

- A vector is also a matrix with 1 column, and its transpose can also be defined.

$$v^\top = \begin{bmatrix} v_1 & v_2 & \cdots & v_n \end{bmatrix} \in \mathbb{R}^{1,n} \quad (4)$$

## Mathematical Notations (7): Tensors

- **Tensor (Tensor):**
  - In this lecture, the term tensor simply refers to a multi-dimensional array. A vector can be regarded as a 1st-order tensor, and a matrix as a 2nd-order tensor. Tensors of 3rd-order or higher are denoted by underlined bold italic uppercase letters, like *A*.
  - Students who have already learned about abstract tensors in mathematics or physics may feel resistant to calling a mere multi-dimensional array a tensor. If one considers that the basis is always fixed to the standard basis and the mathematical tensor is identified with its component representation (which becomes a multi-dimensional array), then there is (arguably) consistency in the terminology.

# **The Idea of PEFT: Representing Model Differences with Low-Dimensional Parameters**

---

## Review of PEFT Motivation

Let's review the motivation for PEFT. We assume a situation where we already have a model  $f_\theta$  that provides a good input-output relationship, and we expect to obtain a better input-output relationship by modifying it slightly. In this situation, the idea of PEFT is to try to represent the part corresponding to that slight modification with a small number of parameters (or a low-dimensional parameter vector).

## Remark on “Slight Modification”

### Remark

Here, “slight modification” is not something to be defined strictly mathematically. Although there are many mathematical criteria for measuring the closeness of functions, it does not necessarily mean “slight” in that sense here. “Slight” here, in the end, can only be defined as a modification that can be represented by a small number of parameters, which becomes a tautology. To be more precise, it is more accurate to say that methods have been empirically found that can achieve good performance improvements with specific modifications represented by a small number of parameters, and these are called PEFT.

## PEFT Adapter: Setup

We will strictly define PEFT as a function. Here, we let the parameter space of the original model  $f_\theta$  be

$$\boldsymbol{\theta} \in \mathbb{R}^{d_{\text{Original}}} \quad (5)$$

and the input and output spaces be  $\mathcal{X}$  and  $\mathcal{Y}$ , respectively.

## Definition: PEFT Adapter

### Definition (PEFT Adapter)

Let  $\mathcal{X}$  be the input space and  $\mathcal{Y}$  be the output space. Fix  $d_{\text{Original}} \in \mathbb{Z}_{>0}$ , and suppose we are given a parametric function  $f(\cdot)$  with parameters on the parameter vector space  $\mathbb{R}^{d_{\text{Original}}}$ . Given  $\theta \in \mathbb{R}^{d_{\text{Original}}}$ , a specific function  $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$  is determined.

## Definition: PEFT Adapter

### Definition (PEFT Adapter)

Let  $\mathcal{X}$  be the input space and  $\mathcal{Y}$  be the output space. Fix  $d_{\text{Original}} \in \mathbb{Z}_{>0}$ , and suppose we are given a parametric function  $f(\cdot)$  with parameters on the parameter vector space  $\mathbb{R}^{d_{\text{Original}}}$ . Given  $\theta \in \mathbb{R}^{d_{\text{Original}}}$ , a specific function  $f_\theta : \mathcal{X} \rightarrow \mathcal{Y}$  is determined.

Fix  $d_{\text{PEFT}} \in \mathbb{Z}_{>0}$ , and let the **PEFT parameter** be  $\psi \in \mathbb{R}^{d_{\text{PEFT}}}$ . At this time, a **PEFT adapter** is a map

$$\text{PEFTAdapter}_{f(\cdot), \psi} : \mathbb{R}^{d_{\text{Original}}} \rightarrow \mathcal{Y}^{\mathcal{X}}. \quad (6)$$

The original model  $\theta \in \mathbb{R}^{d_{\text{Original}}}$  is transformed into  $\text{PEFTAdapter}_{f(\cdot), \psi}(\theta)$ , depending on the small number of parameters  $\psi$ . The specific way this transformation is determined depends on the PEFT method.

## Remark: PEFT Adapter Intuition

### Remark

Definition 1 defines PEFT very abstractly, but the essence is as follows. We fix a base model  $f_\theta$  and construct a new function  $\text{PEFTAdapter}_{f(\cdot), \psi}(\theta)$  through the PEFTAdapter. At this time,  $\text{PEFTAdapter}_{f(\cdot), \psi}(\theta)$  is controlled by the low-dimensional parameter  $\psi$ . That is, it is a framework that allows access to various models in the vicinity of  $f_\theta$  by changing  $\psi$ .

## PEFT as Learning Problem

Under Definition 1, **PEFT is a framework for determining the PEFT parameter  $\psi$  in the PEFT adapter by learning, so that it becomes a “good” function for the desired task.** More specifically, the learning in PEFT involves defining some objective function (e.g., a loss function) as a function of  $\psi$  and minimizing it.

## Motivation for a Simple Example

Since Definition 1 is abstract, let's consider the simplest specific example.  
Consider the case where the original parameter vector  $\theta$  is explicitly divided into a part that is not updated during learning and a part that is updated.

## Example: Simple PEFT by Modifying Only Some Parameters

### Example (Simple PEFT by Modifying Only Some Parameters)

For  $d_{\text{fixed}}, d_{\text{tuned}} \in \mathbb{Z}_{>0}$ , let  $d_{\text{Original}} := d_{\text{fixed}} + d_{\text{tuned}}$  and represent the parameter vector by block partitioning as

$$\boldsymbol{\theta} = \begin{bmatrix} \boldsymbol{\theta}_{\text{fixed}} \\ \boldsymbol{\theta}_{\text{tuned}} \end{bmatrix} \in \mathbb{R}^{d_{\text{Original}}}. \quad (7)$$

Here,  $\boldsymbol{\theta}_{\text{fixed}} \in \mathbb{R}^{d_{\text{fixed}}}$  are the parameters not changed during learning (fixed parameters), and  $\boldsymbol{\theta}_{\text{tuned}} \in \mathbb{R}^{d_{\text{tuned}}}$  are the parameters changed during learning.

## Example: Simple PEFT by Modifying Only Some Parameters

### Example (Simple PEFT by Modifying Only Some Parameters)

For  $d_{\text{fixed}}, d_{\text{tuned}} \in \mathbb{Z}_{>0}$ , let  $d_{\text{Original}} := d_{\text{fixed}} + d_{\text{tuned}}$  and represent the parameter vector by block partitioning as

$$\boldsymbol{\theta} = \begin{bmatrix} \boldsymbol{\theta}_{\text{fixed}} \\ \boldsymbol{\theta}_{\text{tuned}} \end{bmatrix} \in \mathbb{R}^{d_{\text{Original}}}. \quad (7)$$

Here,  $\boldsymbol{\theta}_{\text{fixed}} \in \mathbb{R}^{d_{\text{fixed}}}$  are the parameters not changed during learning (fixed parameters), and  $\boldsymbol{\theta}_{\text{tuned}} \in \mathbb{R}^{d_{\text{tuned}}}$  are the parameters changed during learning.

At this time, we define the SimplePEFTAdapter $_{f(\cdot), \Delta\boldsymbol{\theta}_{\text{tuned}}}$  as follows:

$$\text{SimplePEFTAdapter}_{f(\cdot), \Delta\boldsymbol{\theta}_{\text{tuned}}} \left( \begin{bmatrix} \boldsymbol{\theta}_{\text{fixed}} \\ \boldsymbol{\theta}_{\text{tuned}} \end{bmatrix} \right) := f \begin{bmatrix} \boldsymbol{\theta}_{\text{fixed}} \\ \boldsymbol{\theta}_{\text{tuned}} + \Delta\boldsymbol{\theta}_{\text{tuned}} \end{bmatrix}. \quad (8)$$

## Simple example note

In the previous example,  $\Delta\theta_{\text{tuned}} \in \mathbb{R}^{d_{\text{tuned}}}$  is the difference parameter determined by learning, and it corresponds to the PEFT parameter in this example. That is, matching the notation of Definition 1, we have

$$\psi = \Delta\theta_{\text{tuned}}, \quad d_{\text{PEFT}} = d_{\text{tuned}}. \quad (9)$$

## Remark: Observations from the Simple PEFT Example

### Remark

Let's make some observations from Example 2.

- First, the number of parameters to be determined is  $d_{\text{PEFT}} = d_{\text{tuned}}$ , which can be set freely by the designer. In particular, it can be designed such that  $d_{\text{tuned}} \ll d_{\text{Original}}$ . This is a great advantage in terms of learning stability when the available training data is scarce.
- Next, consider the perspective of memory for storage. If one only wants to obtain the result of PEFT, one may discard the original  $\theta$  and save only

$$\begin{bmatrix} \theta_{\text{fixed}} \\ \theta_{\text{tuned}} + \Delta\theta_{\text{tuned}}. \end{bmatrix} \quad (10)$$

The number of parameters to save in this case is  $d_{\text{Original}}$ , which is the same as saving the original model.

## Remark: Observations from the Simple PEFT Example (2)

### Remark (continued)

- On the other hand, suppose that both the original  $f_\theta$  and the PEFT-modified  $f\begin{bmatrix} \theta_{\text{fixed}} \\ \theta_{\text{tuned}} + \Delta\theta_{\text{tuned}} \end{bmatrix}$  are valuable for applications. In this case, naively thinking, one would save two sets of parameters,  $\theta$  and (10), resulting in storing a total of  $2d_{\text{Original}}$  floating-point values.
- However, in reality, the PEFT-modified (10) can be recovered from the original  $\theta$  and  $\Delta\theta_{\text{tuned}}$  as  $\begin{bmatrix} \theta_{\text{fixed}} \\ \theta_{\text{tuned}} + \Delta\theta_{\text{tuned}} \end{bmatrix} = \begin{bmatrix} \theta_{\text{fixed}} \\ \theta_{\text{tuned}} \end{bmatrix} + \begin{bmatrix} 0 \\ \Delta\theta_{\text{tuned}} \end{bmatrix}$ , so the number of floating-point values that need to be saved is  $d_{\text{Original}} + d_{\text{PEFT}} = d_{\text{Original}} + d_{\text{tuned}}$ . Usually,  $d_{\text{tuned}} \ll d_{\text{Original}}$ , so PEFT can significantly reduce the number of floating-point values to be saved.

## Remark: Observations from the Simple PEFT Example (3)

### Remark (continued)

- In particular, if PEFT is performed separately for multiple purposes, for example, if  $K$  PEFTs are performed and each PEFT result is valuable for applications. Naive storage would require saving  $(K + 1)d_{\text{Original}}$  floating-point values. On the other hand, when using PEFT, one only needs to save the original  $\theta$  and each task-specific  $\Delta\theta_{\text{tuned}}^{(k)}$  ( $k = 1, \dots, K$ ), so the number of floating-point values to be saved is

$$d_{\text{Original}} + Kd_{\text{PEFT}} = d_{\text{Original}} + Kd_{\text{tuned}}. \quad (11)$$

When  $d_{\text{tuned}} \ll d_{\text{Original}}$  and  $K$  is large, this difference is practically very significant.

## PEFT Advantages

Thus, PEFT is advantageous from both the perspective of “number of parameters to update during learning” and “number of parameters to save”. In the next chapter, we will introduce LoRA, which achieves even higher parameter efficiency by imposing further structural constraints on parameters represented as matrices.

## **The Idea of Low Rank Adaptation (LoRA)**

---

## Motivation for Low-Rank Constraints

In the method described above, if  $d_{\text{PEFT}} \ll d_{\text{Original}}$ , the parameter efficiency of PEFT is good, but a concern is whether it can provide practically good differences when  $d_{\text{PEFT}}$  is small. Of course, mathematically considering the worst case, it's impossible for a small number of parameters to substitute for many parameters. The issue here is not the mathematical worst case, but whether, when considering practical utility in natural or social sciences, a better difference can be represented with a small number of parameters.

## Empirical Motivation for Low-Rank Differences

In fact, for the neural networks of generative AI in practical use, it is empirically known that when a part of the parameters can be regarded as a matrix representing a linear transformation, it is better to limit the difference to be represented by a **low-rank matrix** using fewer parameters, rather than representing the difference directly. This can represent practically useful differences despite the small number of parameters [3, 2]. **LoRA (low-rank adaptation)** is based on this idea, representing the difference for matrix-form parameters as a product of a matrix with few rows (a matrix with a small number of rows) and a matrix with few columns (a matrix with a small number of columns).

## Rank of a Matrix via Factorization

As a basic matter of linear algebra, we will strictly define the rank of a matrix based on matrix factorization.

## Definition: Matrix Rank by Factorization

### Definition (Definition of Matrix Rank by Factorization)

Let  $d_{\text{out}}, d_{\text{in}} \in \mathbb{Z}_{>0}$  and  $\mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$ . Consider the set

$$S(\mathbf{W}) := \left\{ r \in \mathbb{Z}_{\geq 0} \mid \exists \mathbf{B} \in \mathbb{R}^{d_{\text{out}}, r}, \mathbf{A} \in \mathbb{R}^{r, d_{\text{in}}} \text{ exist such that } \mathbf{W} = \mathbf{B}\mathbf{A} \right\}. \quad (12)$$

At this time, the **rank** of  $\mathbf{W}$  is defined as  $\text{rank}(\mathbf{W}) := \min S(\mathbf{W})$ . That is,  $\text{rank}(\mathbf{W})$  is the minimum value of the number of column vectors  $r$  required when representing  $\mathbf{W}$  as

$$\mathbf{W} = \mathbf{B}\mathbf{A}, \quad \mathbf{B} \in \mathbb{R}^{d_{\text{out}}, r}, \quad \mathbf{A} \in \mathbb{R}^{r, d_{\text{in}}}, \quad (13)$$

and it represents the “minimum number of vectors” necessary to represent  $\mathbf{W}$ . In particular, if  $\mathbf{W} = \mathbf{0}$ , then  $\text{rank}(\mathbf{W}) = 0$ .

## Proposition: Equivalent Characterizations of Matrix Rank

### Proposition (Equivalent Characterizations of Matrix Rank)

For  $\text{rank}(\mathbf{W})$  defined in Definition 3, the following hold.

- (1)  $\text{rank}(\mathbf{W}) = \dim \mathcal{C}(\mathbf{W})$ . Here,  $\mathcal{C}(\mathbf{W})$  is the subspace spanned by the column vectors of  $\mathbf{W}$ , called the **column space**.
- (2)  $\text{rank}(\mathbf{W}) = \dim \mathcal{R}(\mathbf{W})$ . Here,  $\mathcal{R}(\mathbf{W})$  is the subspace spanned by the row vectors of  $\mathbf{W}$ , called the **row space**.

In particular, the dimensions of the column space and row space of  $\mathbf{W}$  are always equal.

## Remark: Rank, Column Space, and Parameter Count

### Remark

From Proposition 1,  $\text{rank}(\mathbf{W}) = r$  is equivalent to the dimension of the column space of  $\mathbf{W}$  being  $r$ , i.e., “the number of linearly independent column vectors needed to generate all column vectors of  $\mathbf{W}$  is  $r$ ”. On the other hand, according to Definition 3, when  $\text{rank}(\mathbf{W}) = r$ ,  $\mathbf{W}$  can be expressed using  $\mathbf{B} \in \mathbb{R}^{d_{\text{out}}, r}$  and  $\mathbf{A} \in \mathbb{R}^{r, d_{\text{in}}}$  as in equation (13). In this case, the number of elements in  $\mathbf{B}$  is  $d_{\text{out}}r$ , and the number of elements in  $\mathbf{A}$  is  $rd_{\text{in}}$ , so  $\mathbf{W}$  can be described by specifying a total of  $d_{\text{rank}} := d_{\text{out}}r + rd_{\text{in}}$  scalars. In particular, when  $r \ll \min\{d_{\text{out}}, d_{\text{in}}\}$ , we have  $d_{\text{rank}} \ll d_{\text{out}}d_{\text{in}}$  so a low-rank matrix is a class of matrices that can be described with far fewer parameters compared to handling the original full-rank matrix directly. This is perfectly consistent with the motivation of PEFT, which is to represent models and differences as much as possible with a small number of parameters.

## LoRA and Low-Rank Differences

LoRA improves parameter efficiency by artificially constraining the difference matrix  $\Delta W$  to be represented in this form.

## Definition: Difference Matrix in LoRA

### Definition (Difference Matrix in LoRA)

Fix  $d_{\text{out}}, d_{\text{in}} \in \mathbb{Z}_{>0}$ , and let the original matrix parameter be  $\mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$ . Let  $r \in \mathbb{Z}_{>0}$  be a parameter called **rank**, and  $\alpha \in \mathbb{R}$  be a parameter called **scaling factor**.

## Definition: Difference Matrix in LoRA

### Definition (Difference Matrix in LoRA)

Fix  $d_{\text{out}}, d_{\text{in}} \in \mathbb{Z}_{>0}$ , and let the original matrix parameter be  $\mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$ . Let  $r \in \mathbb{Z}_{>0}$  be a parameter called **rank**, and  $\alpha \in \mathbb{R}$  be a parameter called **scaling factor**.

At this time, the difference matrix  $\Delta \mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$  by LoRA is defined as

$$\Delta \mathbf{W} := \frac{\alpha}{r} \mathbf{B} \mathbf{A}. \quad (14)$$

Here,  $\mathbf{A} \in \mathbb{R}^{r, d_{\text{in}}}$ ,  $\mathbf{B} \in \mathbb{R}^{d_{\text{out}}, r}$  are matrices determined by learning, and are the main learning targets of LoRA.

## Definition: Difference Matrix in LoRA

### Definition (Difference Matrix in LoRA)

Fix  $d_{\text{out}}, d_{\text{in}} \in \mathbb{Z}_{>0}$ , and let the original matrix parameter be  $\mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$ . Let  $r \in \mathbb{Z}_{>0}$  be a parameter called **rank**, and  $\alpha \in \mathbb{R}$  be a parameter called **scaling factor**.

At this time, the difference matrix  $\Delta \mathbf{W} \in \mathbb{R}^{d_{\text{out}}, d_{\text{in}}}$  by LoRA is defined as

$$\Delta \mathbf{W} := \frac{\alpha}{r} \mathbf{B} \mathbf{A}. \quad (14)$$

Here,  $\mathbf{A} \in \mathbb{R}^{r, d_{\text{in}}}$ ,  $\mathbf{B} \in \mathbb{R}^{d_{\text{out}}, r}$  are matrices determined by learning, and are the main learning targets of LoRA.

The matrix parameter after applying LoRA is defined as

$$\widetilde{\mathbf{W}} := \mathbf{W} + \Delta \mathbf{W} = \mathbf{W} + \frac{\alpha}{r} \mathbf{B} \mathbf{A}. \quad (15)$$

## Remark on $\alpha$ and $r$ in LoRA

### Remark

As an implementation note regarding LoRA, the  $\alpha/r$  appearing in equation (14) is redundant in a strict sense. That is, the same expressive power can be maintained by setting  $\alpha$  to 1 and incorporating the scalar into  $A$  or  $B$ .

## Remark on $\alpha$ and $r$ in LoRA

### Remark

As an implementation note regarding LoRA, the  $\alpha/r$  appearing in equation (14) is redundant in a strict sense. That is, the same expressive power can be maintained by setting  $\alpha$  to 1 and incorporating the scalar into  $A$  or  $B$ .

In reality, the LoRA paper [3], after including  $\alpha$  and  $r$  in the formulation as above, explicitly states that  $\alpha$  is set equal to the rank  $r$ , i.e.,  $\alpha = r$ , as an experimental setup. In this case, the coefficient  $\alpha/r$  in equation (14) is always 1, which is mathematically equivalent to adding  $\Delta W = BA$ . The inclusion of  $\alpha$  and  $r$  in the formulation is for the convenience of theoretical explanation.

## Remark on $\alpha$ and $r$ in LoRA

### Remark

As an implementation note regarding LoRA, the  $\alpha/r$  appearing in equation (14) is redundant in a strict sense. That is, the same expressive power can be maintained by setting  $\alpha$  to 1 and incorporating the scalar into  $A$  or  $B$ .

In reality, the LoRA paper [3], after including  $\alpha$  and  $r$  in the formulation as above, explicitly states that  $\alpha$  is set equal to the rank  $r$ , i.e.,  $\alpha = r$ , as an experimental setup. In this case, the coefficient  $\alpha/r$  in equation (14) is always 1, which is mathematically equivalent to adding  $\Delta W = BA$ . The inclusion of  $\alpha$  and  $r$  in the formulation is for the convenience of theoretical explanation.

The above setting is in line with the LoRA training script

`train_text_to_image_lora.py` for Stable Diffusion in Hugging Face diffusers.

## Parameter Count in LoRA

Let's count the number of parameters in the difference matrix. Whereas the original full degrees of freedom are  $d_{\text{out}}d_{\text{in}}$ , in LoRA, the degrees of freedom for  $\mathbf{A}, \mathbf{B}$  are

$$d_{\text{LoRA}} := rd_{\text{in}} + d_{\text{out}}r = r(d_{\text{in}} + d_{\text{out}}) \quad (16)$$

and one for the scalar  $\alpha$ , totaling  $r(d_{\text{in}} + d_{\text{out}}) + 1$ . Usually, we choose  $r \ll \min\{d_{\text{in}}, d_{\text{out}}\}$ , so  $d_{\text{LoRA}} \ll d_{\text{out}}d_{\text{in}}$ .

## Example: Comparison of Specific Parameter Counts

### Example (Comparison of Specific Parameter Counts)

As specific numbers, consider the case  $d_{\text{in}} = 768$ ,  $d_{\text{out}} = 3072$ ,  $r = 32$ . This is a size typically appearing in the linear transformations of intermediate layers in Transformer-based models [7, 5].

## Example: Comparison of Specific Parameter Counts

### Example (Comparison of Specific Parameter Counts)

As specific numbers, consider the case  $d_{\text{in}} = 768$ ,  $d_{\text{out}} = 3072$ ,  $r = 32$ . This is a size typically appearing in the linear transformations of intermediate layers in Transformer-based models [7, 5].

First, the number of elements in the original matrix parameter  $\mathbf{W} \in \mathbb{R}^{3072, 768}$  is

$$d_{\text{full}} = d_{\text{out}} d_{\text{in}} = 3072 \times 768 \quad (17)$$

. Calculating this gives

$$d_{\text{full}} = 3072 \times 768 = 2,359,296 \quad (18)$$

scalars.

## Example: Comparison of Specific Parameter Counts (2)

### Example (continued)

On the other hand, the number of elements in  $A, B$  representing the difference by LoRA is

$$d_{\text{LoRA}} = r(d_{\text{in}} + d_{\text{out}}) = 32(768 + 3072) = 32 \times 3840 = 122,880, \quad (19)$$

and even adding the scalar  $\alpha$ , it is 122,881.

## Example: Comparison of Specific Parameter Counts (2)

### Example (continued)

On the other hand, the number of elements in  $A, B$  representing the difference by LoRA is

$$d_{\text{LoRA}} = r(d_{\text{in}} + d_{\text{out}}) = 32(768 + 3072) = 32 \times 3840 = 122,880, \quad (19)$$

and even adding the scalar  $\alpha$ , it is 122,881.

Therefore, compared to representing the difference matrix at full size, the number of parameters required is

$$\frac{d_{\text{LoRA}}}{d_{\text{full}}} \approx \frac{1.23 \times 10^5}{2.36 \times 10^6} \approx 0.052, \quad (20)$$

which is reduced to about 1/19.

## Setup: LoRA for General Neural Networks

Next, we define the framework for applying LoRA to a general neural network containing multiple matrix parameters.

## Definition: LoRA for General Neural Networks

### Definition (LoRA for General Neural Networks)

Let  $f_{W_1, W_2, \dots, W_K, \theta_{\text{others}}}$  be a neural network consisting of  $K$  matrix parameters and other parameters. That is,  $W_k \in \mathbb{R}^{d_{\text{out},k}, d_{\text{in},k}}$  ( $k = 1, \dots, K$ ) and

$$\theta_{\text{others}} \in \mathbb{R}^{d_{\text{others}}}. \quad (21)$$

Let the entire set of parameters be

$$\Theta = (W_1, \dots, W_K, \theta_{\text{others}}), \quad (22)$$

and this is the parameter of the model  $f$ .

## Definition: LoRA for General Neural Networks (2)

### Definition (LoRA for General Neural Networks, continued)

For each  $k \in [1, K]_{\mathbb{Z}}$ , fix a rank  $r_k \in \mathbb{Z}_{>0}$ , and let  $\mathbf{A}_k \in \mathbb{R}^{r_k, d_{\text{in},k}}$ ,  $\mathbf{B}_k \in \mathbb{R}^{d_{\text{out},k}, r_k}$  be the LoRA matrix parameters, and  $\alpha_k \in \mathbb{R}$  be the LoRA scaling factor. Also, let  $\eta \in \mathbb{R}$  be a parameter that controls the overall scaling of LoRA.  $\eta$  is usually a hyperparameter set at inference time and is often not a learning target.

## Definition: LoRA for General Neural Networks (2)

### Definition (LoRA for General Neural Networks, continued)

For each  $k \in [1, K]_{\mathbb{Z}}$ , fix a rank  $r_k \in \mathbb{Z}_{>0}$ , and let  $\mathbf{A}_k \in \mathbb{R}^{r_k, d_{\text{in},k}}$ ,  $\mathbf{B}_k \in \mathbb{R}^{d_{\text{out},k}, r_k}$  be the LoRA matrix parameters, and  $\alpha_k \in \mathbb{R}$  be the LoRA scaling factor. Also, let  $\eta \in \mathbb{R}$  be a parameter that controls the overall scaling of LoRA.  $\eta$  is usually a hyperparameter set at inference time and is often not a learning target.

At this time, the **LoRA adapter** is defined as

$$\begin{aligned} & \text{LoRA}_{f(\cdot), \eta, \alpha_1, \mathbf{A}_1, \mathbf{B}_1, \dots, \alpha_K, \mathbf{A}_K, \mathbf{B}_K}(\mathbf{W}_1, \dots, \mathbf{W}_K, \theta_{\text{others}}) \\ &:= f_{\mathbf{W}_1 + \eta \frac{\alpha_1}{r_1} \mathbf{B}_1 \mathbf{A}_1, \mathbf{W}_2 + \eta \frac{\alpha_2}{r_2} \mathbf{B}_2 \mathbf{A}_2, \dots, \mathbf{W}_K + \eta \frac{\alpha_K}{r_K} \mathbf{B}_K \mathbf{A}_K, \theta_{\text{others}}} \quad (23) \end{aligned}$$

Here, for  $k = 1, 2, \dots, K$ ,  $r_k$  is a predetermined positive integer and is not changed during learning. Also,  $\alpha_k$  is usually a fixed positive number.

## Remark: Controlling LoRA Strength

### Remark

In Definition 7,  $\eta$  is a parameter that adjusts the “strength” of LoRA. In the learning phase, learning is often performed with  $\eta = 1$ , and by changing  $\eta$  in the inference phase,

$$\widetilde{\mathbf{W}}_k(\eta) = \mathbf{W}_k + \eta \frac{\alpha_k}{r_k} \mathbf{B}_k \mathbf{A}_k \quad (24)$$

, the contribution of the modification by LoRA can be strengthened or weakened. Setting  $\eta = 0$  returns to the original model  $f_{\mathbf{W}_1, \dots, \mathbf{W}_K, \theta_{\text{others}}}$ .

## LoRA Training: Objective Function

Under Definition 7, training LoRA means determining  $\mathbf{A}_k, \mathbf{B}_k$  ( $k = 1, \dots, K$ ) based on the training data. Here, we assume that the original matrix parameters  $\mathbf{W}_k$  and other parameters  $\theta_{\text{others}}$  are fixed and not updated.

## LoRA Training: Objective Function

Under Definition 7, training LoRA means determining  $\mathbf{A}_k, \mathbf{B}_k$  ( $k = 1, \dots, K$ ) based on the training data. Here, we assume that the original matrix parameters  $\mathbf{W}_k$  and other parameters  $\theta_{\text{others}}$  are fixed and not updated.

Given  $m$  input-output pairs  $(\mathbf{x}^{(i)}, \mathbf{y}^{(i)})$  for  $i = 1, \dots, m$ , define the objective function (loss function) as

$$\mathcal{L}(\alpha_1, \mathbf{A}_1, \mathbf{B}_1, \dots, \alpha_K, \mathbf{A}_K, \mathbf{B}_K) := \frac{1}{m} \sum_{i=1}^m \ell(\mathbf{y}^{(i)}, f^{\text{LoRA}}(\mathbf{x}^{(i)})) \quad (25)$$

## LoRA Training: Model with LoRA

Here,

$$f^{\text{LoRA}}(\mathbf{x}) \coloneqq \text{LoRA}_{f(\cdot), \eta=1, \alpha_1, \mathbf{A}_1, \mathbf{B}_1, \dots, \alpha_K, \mathbf{A}_K, \mathbf{B}_K}(\mathbf{W}_1, \dots, \mathbf{W}_K, \theta_{\text{others}})(\mathbf{x}) \quad (26)$$

$$= f_{\mathbf{W}_1 + \frac{\alpha_1}{r_1} \mathbf{B}_1 \mathbf{A}_1, \dots, \mathbf{W}_K + \frac{\alpha_K}{r_K} \mathbf{B}_K \mathbf{A}_K, \theta_{\text{others}}}(\mathbf{x}) \quad (27)$$

, and  $\ell$  is a task-dependent loss function, such as mean squared error or cross-entropy.

## Definition: LoRA Training Problem

### Definition (LoRA Training Problem)

The LoRA training problem is, given the fixed original parameters  $\mathbf{W}_1, \dots, \mathbf{W}_K, \theta_{\text{others}}$ , to solve

$$\underset{\mathbf{A}_1, \mathbf{B}_1, \dots, \mathbf{A}_K, \mathbf{B}_K}{\text{Minimize}} \quad \mathcal{L}(\mathbf{A}_1, \mathbf{B}_1, \dots, \mathbf{A}_K, \mathbf{B}_K) \quad (28)$$

## Remark: Gradient-Based Training for LoRA

### Remark

In actual training, a gradient method (stochastic gradient descent, AdamW, etc.) is used to numerically solve the minimization problem in Definition 9. That is, at each step,

$$\frac{\partial \mathcal{L}}{\partial \mathbf{A}_k}, \quad \frac{\partial \mathcal{L}}{\partial \mathbf{B}_k} \tag{29}$$

are obtained by backpropagation, and  $\mathbf{A}_k, \mathbf{B}_k$  are updated based on them. At this time, gradients with respect to  $\mathbf{W}_k$  and  $\theta_{\text{others}}$  are calculated but not used for updates (they are not updated). This allows training only the LoRA parameters without changing the original model.

## LoRA Strength Control at Inference Time

After the LoRA parameters  $\alpha_k, \mathbf{A}_k, \mathbf{B}_k$  are determined by Definition 9, the strength of the LoRA contribution can be controlled using  $\eta$  from Definition 7 in the inference phase.

## LoRA Strength Control at Inference Time

After the LoRA parameters  $\alpha_k, \mathbf{A}_k, \mathbf{B}_k$  are determined by Definition 9, the strength of the LoRA contribution can be controlled using  $\eta$  from Definition 7 in the inference phase.

That is, at inference time, by specifying an arbitrary  $\eta \in \mathbb{R}$  and defining

$$f^{\text{LoRA}, \eta}(\mathbf{x}) = f_{\mathbf{W}_1 + \eta \frac{\alpha_1}{r_1} \mathbf{B}_1 \mathbf{A}_1, \dots, \mathbf{W}_K + \eta \frac{\alpha_K}{r_K} \mathbf{B}_K \mathbf{A}_K, \theta_{\text{others}}}(\mathbf{x}) \quad (30)$$

, one can return to the original model with  $\eta = 0$ , use the same strength as during training with  $\eta = 1$ , or strengthen or weaken the LoRA contribution more than usual with  $|\eta| > 1$ .

## LoRA Strength Control at Inference Time

After the LoRA parameters  $\alpha_k, \mathbf{A}_k, \mathbf{B}_k$  are determined by Definition 9, the strength of the LoRA contribution can be controlled using  $\eta$  from Definition 7 in the inference phase.

That is, at inference time, by specifying an arbitrary  $\eta \in \mathbb{R}$  and defining

$$f^{\text{LoRA}, \eta}(\mathbf{x}) = f_{\mathbf{W}_1 + \eta \frac{\alpha_1}{r_1} \mathbf{B}_1 \mathbf{A}_1, \dots, \mathbf{W}_K + \eta \frac{\alpha_K}{r_K} \mathbf{B}_K \mathbf{A}_K, \theta_{\text{others}}}(\mathbf{x}) \quad (30)$$

, one can return to the original model with  $\eta = 0$ , use the same strength as during training with  $\eta = 1$ , or strengthen or weaken the LoRA contribution more than usual with  $|\eta| > 1$ .

In implementation, by allowing the user to explicitly specify  $\eta$ , the “strength of the LoRA flavor” can be easily adjusted. This is frequently used for purposes such as adjusting style transfer or the tendency of prompt responses.

## **Specific Application of LoRA: Example with Stable Diffusion 1.5**

---

## Components of Stable Diffusion 1.5

In this section, we will outline the application of LoRA in Stable Diffusion 1.5 [6, 1] as a specific example. Here, we will mathematically describe which matrix parameters LoRA is applied to, referring to the layer names in the Python libraries diffusers<sup>1</sup> and transformers<sup>2</sup>.

---

<sup>1</sup>Official repository for the diffusers library: <https://github.com/huggingface/diffusers>

## Components of Stable Diffusion 1.5

In this section, we will outline the application of LoRA in Stable Diffusion 1.5 [6, 1] as a specific example. Here, we will mathematically describe which matrix parameters LoRA is applied to, referring to the layer names in the Python libraries diffusers<sup>1</sup> and transformers<sup>2</sup>.

Stable Diffusion 1.5 is broadly composed of the following components:

- **Text encoder**: Usually a CLIPTextModel from transformers or a similar model, which generates embeddings (condition vectors) from the input text token sequence [4].
- **U-Net denoiser** (U-Net style noise estimator): Implemented as UNet2DConditionModel in diffusers, it estimates noise from a noisy latent variable. It includes self-attention and cross-attention blocks.
- **VAE (variational auto-encoder)**: Implemented as AutoencoderKL or similar, it handles the mutual conversion between images and latent representations.

<sup>1</sup>Official repository for the diffusers library: <https://github.com/huggingface/diffusers>

## LoRA Application Points in Stable Diffusion 1.5

In practice, the main places where LoRA is applied are the following two:

- Targeting only the U-Net: Apply LoRA to the linear layers related to self-attention and cross-attention in the noise estimator (especially `to_q`, `to_k`, `to_v`, `to_out`).
- Also applying to the text encoder: In addition to the above, apply LoRA to the self-attention layers (`self.attn.q_proj`, `k_proj`, `v_proj`, `out_proj`) or MLP parts of the CLIPTextModel.

## LoRA Application Points in Stable Diffusion 1.5

In practice, the main places where LoRA is applied are the following two:

- Targeting only the U-Net: Apply LoRA to the linear layers related to self-attention and cross-attention in the noise estimator (especially `to_q`, `to_k`, `to_v`, `to_out`).
- Also applying to the text encoder: In addition to the above, apply LoRA to the self-attention layers (`self.attn.q_proj`, `k_proj`, `v_proj`, `out_proj`) or MLP parts of the CLIPTextModel.

In the following, we will define the functions represented by each layer, separating input variables and parameters, and strictly describe in mathematical formulas which matrices LoRA is applied to.

## Dimensions in CrossAttention

Let  $d_{\text{model}}$  be the dimension of the feature vectors,  $d_{\text{head}}$  be the dimension of each head, and  $n_{\text{head}}$  be the number of heads, such that

$$d_{\text{model}} = n_{\text{head}} d_{\text{head}} \quad (31)$$

## Definition: CrossAttention Layer in U-Net i

Let the query input sequence and the key/value input sequences be

$$\mathbf{X}_q \in \mathbb{R}^{L_q, d_{\text{model}}}, \quad \mathbf{X}_k \in \mathbb{R}^{L_k, d_{\text{model}}}, \quad \mathbf{X}_v \in \mathbb{R}^{L_k, d_{\text{model}}} \quad (32)$$

. Here,  $L_q, L_k$  are the query sequence length and the key/value sequence length, respectively.

## Definition: CrossAttention Layer in U-Net ii

The CrossAttention layer of the U-Net has the following matrix parameters:

$$\mathbf{W}_q \in \mathbb{R}^{d_{\text{model}}, d_{\text{model}}}, \quad (33)$$

$$\mathbf{W}_k \in \mathbb{R}^{d_{\text{model}}, d_{\text{model}}}, \quad (34)$$

$$\mathbf{W}_v \in \mathbb{R}^{d_{\text{model}}, d_{\text{model}}}, \quad (35)$$

$$\mathbf{W}_o \in \mathbb{R}^{d_{\text{model}}, d_{\text{model}}}. \quad (36)$$

These correspond to `to_q`, `to_k`, `to_v`, `to_out` in the diffusers implementation.

## Definition: CrossAttention Layer in U-Net iii

First, by linear transformation,

$$\mathbf{Q} := \mathbf{X}_q \mathbf{W}_q^\top \in \mathbb{R}^{L_q, d_{\text{model}}}, \quad (37)$$

$$\mathbf{K} := \mathbf{X}_k \mathbf{W}_k^\top \in \mathbb{R}^{L_k, d_{\text{model}}}, \quad (38)$$

$$\mathbf{V} := \mathbf{X}_v \mathbf{W}_v^\top \in \mathbb{R}^{L_k, d_{\text{model}}} \quad (39)$$

are determined.  $\mathbf{Q}, \mathbf{K}, \mathbf{V}$  are split for each head, and written as

$$\mathbf{Q} \leftrightarrow (\mathbf{Q}^{(h)})_{h=1}^{n_{\text{head}}}, \quad \mathbf{Q}^{(h)} \in \mathbb{R}^{L_q, d_{\text{head}}}, \quad (40)$$

$$\mathbf{K} \leftrightarrow (\mathbf{K}^{(h)})_{h=1}^{n_{\text{head}}}, \quad \mathbf{K}^{(h)} \in \mathbb{R}^{L_k, d_{\text{head}}}, \quad (41)$$

$$\mathbf{V} \leftrightarrow (\mathbf{V}^{(h)})_{h=1}^{n_{\text{head}}}, \quad \mathbf{V}^{(h)} \in \mathbb{R}^{L_k, d_{\text{head}}} \quad (42)$$

## Definition: CrossAttention Layer in U-Net iv

Scaled dot-product attention is calculated for each head. That is,

$$\mathbf{S}^{(h)} := \frac{1}{\sqrt{d_{\text{head}}}} \mathbf{Q}^{(h)} \mathbf{K}^{(h)\top} \in \mathbb{R}^{L_q, L_k} \quad (43)$$

, and applying softmax row-wise,

$$\mathbf{A}^{(h)} := \text{softmax}_{\text{row}}(\mathbf{S}^{(h)}) \in \mathbb{R}^{L_q, L_k} \quad (44)$$

is determined. Then

$$\mathbf{O}^{(h)} := \mathbf{A}^{(h)} \mathbf{V}^{(h)} \in \mathbb{R}^{L_q, d_{\text{head}}} \quad (45)$$

is obtained.

## Definition: CrossAttention Layer in U-Net v

The outputs for  $h = 1, \dots, n_{\text{head}}$  are concatenated in the feature dimension direction

$$\mathbf{O} := \text{Concat}_{h=1}^{n_{\text{head}}} \mathbf{O}^{(h)} \in \mathbb{R}^{L_q, d_{\text{model}}} \quad (46)$$

, and finally

$$\mathbf{Y} := \mathbf{O} \mathbf{W}_o^\top \in \mathbb{R}^{L_q, d_{\text{model}}} \quad (47)$$

is the output of the CrossAttention layer.

## Remark: CrossAttention Intuition

### Remark

Definition 51 writes down the standard multi-head attention of the Transformer [7], adapted for the attention blocks of the U-Net. Although batch dimensions, positional encodings, etc., are included in the implementation, the important point is that the linear layers (37)–(39), (47) targeted by LoRA have been explicitly shown in mathematical formulas.

## Applying LoRA to CrossAttention

A typical method for applying LoRA is to add separate LoRAs to each of  $\mathbf{W}_q$ ,  $\mathbf{W}_k$ ,  $\mathbf{W}_v$ ,  $\mathbf{W}_o$ . For example, consider LoRA for  $\mathbf{W}_q$ .

## Applying LoRA to CrossAttention

A typical method for applying LoRA is to add separate LoRAs to each of  $\mathbf{W}_q, \mathbf{W}_k, \mathbf{W}_v, \mathbf{W}_o$ . For example, consider LoRA for  $\mathbf{W}_q$ .

Prepare rank  $r_q$ , scaling factor  $\alpha_q$ , matrices  $\mathbf{A}_q \in \mathbb{R}^{r_q, d_{\text{model}}}$ ,  $\mathbf{B}_q \in \mathbb{R}^{d_{\text{model}}, r_q}$ , and

$$\widetilde{\mathbf{W}}_q := \mathbf{W}_q + \eta \frac{\alpha_q}{r_q} \mathbf{B}_q \mathbf{A}_q \quad (48)$$

is determined. Similarly,

$$\begin{aligned} \widetilde{\mathbf{W}}_k &:= \mathbf{W}_k + \eta \frac{\alpha_k}{r_k} \mathbf{B}_k \mathbf{A}_k, & \widetilde{\mathbf{W}}_v &:= \mathbf{W}_v + \eta \frac{\alpha_v}{r_v} \mathbf{B}_v \mathbf{A}_v, & \widetilde{\mathbf{W}}_o &:= \mathbf{W}_o + \eta \frac{\alpha_o}{r_o} \mathbf{B}_o \mathbf{A}_o \end{aligned} \quad (49)$$

are defined. Here,  $r_q, r_k, r_v, r_o$  are the ranks for each respective matrix, and in implementation, a common value is often used.

## Applying LoRA to CrossAttention

A typical method for applying LoRA is to add separate LoRAs to each of  $\mathbf{W}_q, \mathbf{W}_k, \mathbf{W}_v, \mathbf{W}_o$ . For example, consider LoRA for  $\mathbf{W}_q$ .

Prepare rank  $r_q$ , scaling factor  $\alpha_q$ , matrices  $\mathbf{A}_q \in \mathbb{R}^{r_q, d_{\text{model}}}$ ,  $\mathbf{B}_q \in \mathbb{R}^{d_{\text{model}}, r_q}$ , and

$$\widetilde{\mathbf{W}}_q := \mathbf{W}_q + \eta \frac{\alpha_q}{r_q} \mathbf{B}_q \mathbf{A}_q \quad (48)$$

is determined. Similarly,

$$\begin{aligned} \widetilde{\mathbf{W}}_k &:= \mathbf{W}_k + \eta \frac{\alpha_k}{r_k} \mathbf{B}_k \mathbf{A}_k, & \widetilde{\mathbf{W}}_v &:= \mathbf{W}_v + \eta \frac{\alpha_v}{r_v} \mathbf{B}_v \mathbf{A}_v, & \widetilde{\mathbf{W}}_o &:= \mathbf{W}_o + \eta \frac{\alpha_o}{r_o} \mathbf{B}_o \mathbf{A}_o \end{aligned} \quad (49)$$

are defined. Here,  $r_q, r_k, r_v, r_o$  are the ranks for each respective matrix, and in implementation, a common value is often used.

The CrossAttention layer after applying LoRA is defined as in Definition 51, but replacing  $\mathbf{W}_q, \mathbf{W}_k, \mathbf{W}_v, \mathbf{W}_o$  with  $\widetilde{\mathbf{W}}_q, \widetilde{\mathbf{W}}_k, \widetilde{\mathbf{W}}_v, \widetilde{\mathbf{W}}_o$  respectively in (37)–(47).

## Dimensions in CLIP Text Self-Attention

The text encoder of Stable Diffusion 1.5 is the CLIP text model [4], implemented as `CLIPTextModel` in `transformers`. Here too, LoRA is often applied to the linear transformations of the self-attention layers.

## Dimensions in CLIP Text Self-Attention

The text encoder of Stable Diffusion 1.5 is the CLIP text model [4], implemented as `CLIPTextModel` in `transformers`. Here too, LoRA is often applied to the linear transformations of the self-attention layers.

Let  $d_{\text{model}}^{\text{text}}$  be the dimension of the text embedding,  $n_{\text{head}}^{\text{text}}$  be the number of heads, and  $d_{\text{head}}^{\text{text}}$  be the dimension of each head.

$$d_{\text{model}}^{\text{text}} = n_{\text{head}}^{\text{text}} d_{\text{head}}^{\text{text}} \quad (50)$$

. Let the input sequence be

$$\mathbf{X}_{\text{text}} \in \mathbb{R}^{L_{\text{text}}, d_{\text{model}}^{\text{text}}} \quad (51)$$

## Definition: CLIP Text Self-Attention Layer

### Definition (CLIP Text Self-Attention Layer)

The self-attention layer of the CLIP text model has the following matrix parameters:  $\mathbf{W}_q^{\text{text}}, \mathbf{W}_k^{\text{text}}, \mathbf{W}_v^{\text{text}}, \mathbf{W}_o^{\text{text}} \in \mathbb{R}^{d_{\text{model}}^{\text{text}}, d_{\text{model}}^{\text{text}}}$ .

Query, key, and value are defined as

$$\mathbf{Q}^{\text{text}} := \mathbf{X}_{\text{text}} \mathbf{W}_q^{\text{text}\top}, \quad (52)$$

$$\mathbf{K}^{\text{text}} := \mathbf{X}_{\text{text}} \mathbf{W}_k^{\text{text}\top}, \quad (53)$$

$$\mathbf{V}^{\text{text}} := \mathbf{X}_{\text{text}} \mathbf{W}_v^{\text{text}\top}, \quad (54)$$

, and then multi-head self-attention is calculated similarly to Definition 51, and the output is

$$\mathbf{Y}^{\text{text}} := \mathbf{O}^{\text{text}} \mathbf{W}_o^{\text{text}\top}. \quad (55)$$

## LoRA for CLIP Text Self-Attention

The application of LoRA is similar to the CrossAttention in the U-Net, and for each matrix parameter,

$$\tilde{\mathbf{W}}_q^{\text{text}} = \mathbf{W}_q^{\text{text}} + \eta \frac{\alpha_q^{\text{text}}}{r_q^{\text{text}}} \mathbf{B}_q^{\text{text}} \mathbf{A}_q^{\text{text}}, \quad (56)$$

$$\tilde{\mathbf{W}}_k^{\text{text}} = \mathbf{W}_k^{\text{text}} + \eta \frac{\alpha_k^{\text{text}}}{r_k^{\text{text}}} \mathbf{B}_k^{\text{text}} \mathbf{A}_k^{\text{text}}, \quad (57)$$

$$\tilde{\mathbf{W}}_v^{\text{text}} = \mathbf{W}_v^{\text{text}} + \eta \frac{\alpha_v^{\text{text}}}{r_v^{\text{text}}} \mathbf{B}_v^{\text{text}} \mathbf{A}_v^{\text{text}}, \quad (58)$$

$$\tilde{\mathbf{W}}_o^{\text{text}} = \mathbf{W}_o^{\text{text}} + \eta \frac{\alpha_o^{\text{text}}}{r_o^{\text{text}}} \mathbf{B}_o^{\text{text}} \mathbf{A}_o^{\text{text}} \quad (59)$$

are determined. This modifies the self-attention of the text encoder by LoRA. In practice, by combining the LoRA of the text encoder with the LoRA of the U-Net, both the tendency of prompt interpretation and the image generation process are controlled.

## Remark: Fusing LoRA vs. On-the-fly Application i

There are at least the following two options for how to specifically use the LoRA-applied matrix  $\tilde{\mathbf{W}} = \mathbf{W} + \frac{\alpha}{r} \mathbf{B} \mathbf{A}$  in calculations.

- **Pre-fusing method:** Before performing inference,  $\mathbf{W}_{\text{LoRA}} := \mathbf{W} + \frac{\alpha}{r} \mathbf{B} \mathbf{A}$  is calculated once and saved, and thereafter  $\mathbf{W}_{\text{LoRA}}$  is always used to calculate  $\mathbf{W}_{\text{LoRA}} \mathbf{x}$ . In the diffusers library, this is implemented as `fuse_lora()`. This method is advantageous in terms of computation, because once  $\mathbf{W}_{\text{LoRA}}$  is calculated, there is no need to calculate  $\mathbf{B} \mathbf{A} \mathbf{x}$  every time during numerous inferences (image generation steps).

## Remark: Fusing LoRA vs. On-the-fly Application ii

- **Applying LoRA during each forward pass method:** Even at inference time, the original matrix  $W$  is kept as is, and for the input  $x$ ,

$$Wx + \frac{\alpha}{r} B Ax \quad (60)$$

is calculated every time. This method is highly flexible when one wants to switch multiple LoRAs at once or turn LoRA on/off during inference. It is also useful in situations where one absolutely does not want to change the original  $W$ .

## Remark: Fusing LoRA vs. On-the-fly Application iii

In general, the on-the-fly method is convenient in situations where LoRA application and removal are frequently repeated, while the fuse method is advantageous in terms of computational efficiency when a certain LoRA is fixed for a long time to perform a large amount of inference.

## **Summary and Next Lecture Preview**

---

## Summary

We briefly review the Learning Outcomes set forth in this lecture.

- LoRA is a type of parameter-efficient fine-tuning (PEFT) that can efficiently realize a better model for solving a task by determining a small number of parameters that represent the difference from the original model.

## Summary

We briefly review the Learning Outcomes set forth in this lecture.

- LoRA is a type of parameter-efficient fine-tuning (PEFT) that can efficiently realize a better model for solving a task by determining a small number of parameters that represent the difference from the original model.
- LoRA requires changing only a small number of parameters, yet possesses sufficient expressive power for applications. This brings practical advantages such as improved learning stability with small data, reduction in storage usage, and ease of managing numerous task-specific adapters.

## Summary

We briefly review the Learning Outcomes set forth in this lecture.

- LoRA is a type of parameter-efficient fine-tuning (PEFT) that can efficiently realize a better model for solving a task by determining a small number of parameters that represent the difference from the original model.
- LoRA requires changing only a small number of parameters, yet possesses sufficient expressive power for applications. This brings practical advantages such as improved learning stability with small data, reduction in storage usage, and ease of managing numerous task-specific adapters.
- Mathematically, LoRA restricts the difference matrix to one represented by the product of a matrix with few rows and a matrix with few columns, and determines those matrices with few rows and few columns.

## Next Lecture Preview

Next time, we will explain the **relationship between the number of parameters or the size of the function space, and the amount of data required for learning.** That is, we will outline the relationship between model capacity, generalization performance, and the required sample size from the perspective of statistics and learning theory, and consider how methods like PEFT and LoRA can contribute to data efficiency.

## References i

- [1] Stability AI.  
**Stable diffusion v1-5 model card.**  
Technical report, 2022.  
Model card and weights.
- [2] Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer.  
**Qlora: Efficient finetuning of quantized LLMs.**  
Advances in Neural Information Processing Systems, 2023.  
NeurIPS 2023.

## References ii

- [3] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Lu Wang, and Weizhu Chen.  
**Lora: Low-rank adaptation of large language models.**  
arXiv preprint arXiv:2106.09685, 2021.
- [4] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever.  
**Learning transferable visual models from natural language supervision.**  
Proceedings of the 38th International Conference on Machine Learning, pages 8748–8763, 2021.

## References iii

- [5] Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever.

**Language models are unsupervised multitask learners.**

OpenAI Technical Report, 2019.

Technical report.

- [6] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer.

**High-resolution image synthesis with latent diffusion models.**

Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pages 10684–10695, 2022.

- [7] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones,  
Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin.  
**Attention is all you need.**  
Advances in Neural Information Processing Systems, 30, 2017.